

Chapter 09: Recommender Systems

Learning Objectives

- Understand recommender systems and their business applications.
- Learn about the datasets and algorithm required for building recommendation systems.
- Learn recommender system development techniques such as association rules and collaborative filtering.
- Learn how to build and evaluate recommendation systems using Python libraries.

OVERVIEW

Three algorithms that are widely used for building recommendation systems:

- **Association Rules**
- **Collaborative Filtering**
- **Matrix Factorization**

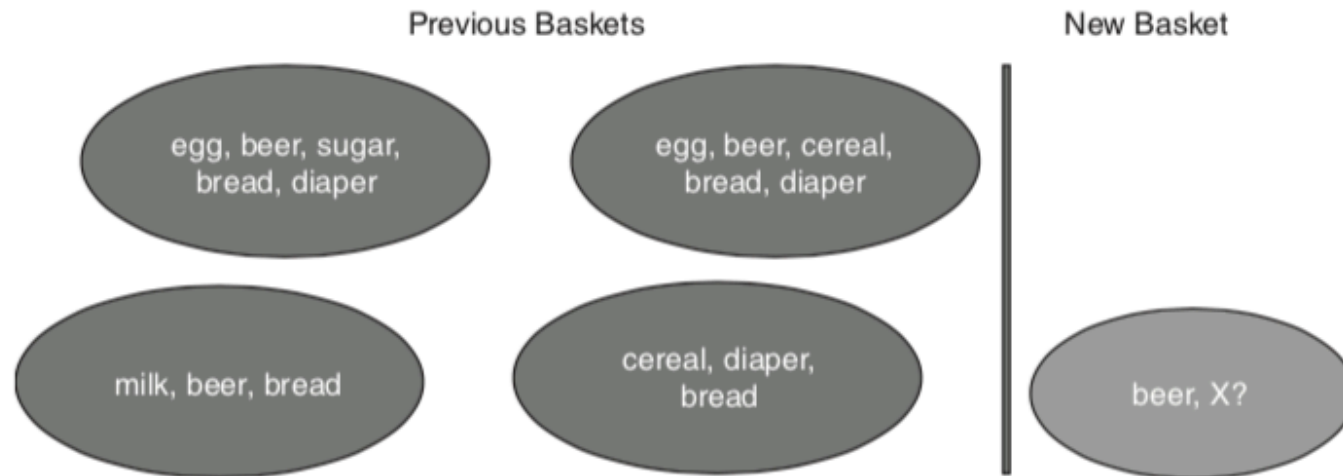
Datasets:

- groceries.csv: contains transactions of a grocery store and can be downloaded from http://www.sci.csueastbay.edu/~esuess/classes/Statistics_6620/Presentations/ml13/groceries.csv.
- Movie Lens: contains ratings and tag applications movies. The dataset can be downloaded from the link <https://grouplens.org/datasets/movielens/>.

ASSOCIATION RULES (ASSOCIATION RULE MINING)

1. Basket 1: egg, beer, sugar, bread, diaper
2. Basket 2: egg, beer, cereal, bread, diaper
3. Basket 3: milk, beer, bread
4. Basket 4: cereal, diaper, bread

$\{\text{diapers}\} \rightarrow \{\text{beer}\}$



METRICS - SUPPORT

Assume that X and Y are items being considered. Let

1. N be the total number of baskets.
2. N_{XY} represent the number of baskets in which X and Y appear together.
3. N_X represent the number of baskets in which X appears.
4. N_Y represent the number of baskets in which Y appears.

Then the support between X and Y , $\text{Support}(X, Y)$, is given by

$$\text{Support}(X, Y) = \frac{N_{XY}}{N}$$

METRICS - CONFIDENCE

Confidence measures the proportion of the transactions that contain X , which also contain Y . X is called antecedent and Y is called consequent. Confidence can be calculated using the following formula:

$$\text{Confidence}(X \rightarrow Y) = P(Y | X) = \frac{N_{XY}}{N_X}$$

where $P(Y|X)$ is the conditional probability of Y given X .

METRICS - LIFT

Lift is calculated using the following formula:

$$\text{Lift} = \frac{\text{Support}(X,Y)}{\text{Support}(X) \times \text{Support}(Y)} = \frac{N_{XY}}{N_X N_Y}$$

Lift can be interpreted as the degree of association between two items.

- Lift value 1 – The items are independent (no association)
- Lift value of less than 1 – The products are substitution (purchase one product will decrease the probability of purchase of the other product)
- Lift value of greater than 1 - purchase of Product X will increase the probability of purchase of Product Y

APPLYING ASSOCIATION RULES

- The code opens the file *groceries.csv*
- Reads all the lines from the file
- Removes leading or trailing white spaces from each line
- Splits each line by a comma to extract items
- Stores the items in each line in a list

```
all_txns[0:5]
```

The output is shown below:

```
[[ 'citrus fruit', 'semi-finished bread', 'margarine', 'ready soups'],  
 [ 'tropical fruit', 'yogurt', 'coffee'],  
 [ 'whole milk'],  
 [ 'pip fruit', 'yogurt', 'cream cheese', 'meat spreads'],  
 [ 'other vegetables',  
  'whole milk',  
  'condensed milk',  
  'long life bakery product']]
```


ENCODING THE TRANSACTIONS

```
import pandas as pd
import numpy as np
from mlxtend.preprocessing import OnehotTransactions
from mlxtend.frequent_patterns import apriori, association_rules

# Initialize OnehotTransactions
one_hot_encoding = OnehotTransactions()

# Transform the data into one-hot-encoding format
one_hot_txns = one_hot_encoding.fit(all_txns).transform(all_txns)

# Convert the matrix into the dataframe.
one_hot_txns_df = pd.DataFrame(one_hot_txns, columns=one_hot_encoding.columns_)
```

	Berries	Beverages	Bottled beer	Bottled water	Brandy	Brown bread	Butter	Butter milk	Cake bar	Candles
5	0	0	0	0	0	0	1	0	0	0
6	0	0	0	0	0	0	0	0	0	0
7	0	0	1	0	0	0	0	0	0	0
8	0	0	0	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0	0	0	0

GENERATING ASSOCIATION RULES

If we have 171 items across all transactions, for itemset containing 2 items in each set, the total number of item sets will be $^{171}C_2$ **14535**

```
frequent_itemsets = apriori(one_hot_txns_df, min_support=0.02, use_colnames=True)
```

	Support	Itemsets
60	0.020437	[bottled beer, whole milk]
52	0.033859	[sugar]
89	0.035892	[other vegetables, tropical fruit]
105	0.021047	[root vegetables, tropical fruit]
88	0.032740	[other vegetables, soda]
16	0.058058	[coffee]
111	0.024504	[shopping bags, whole milk]
36	0.079817	[newspapers]
119	0.056024	[whole milk, yogurt]
55	0.071683	[whipped/sour cream]

ASSOCIATION RULES

```
rules = association_rules(frequent_itemsets, # itemsets  
                          metric="lift", # lift  
                          min_threshold=1 )
```

	Antecedants	Consequents	Support	Confidence	Lift
7	(soda)	(rolls/buns)	0.174377	0.219825	1.195124
55	(yogurt)	(bottled water)	0.139502	0.164723	1.490387
74	(soda)	(yogurt)	0.174377	0.156851	1.124368
89	(root vegetables)	(whole milk)	0.108998	0.448694	1.756031
59	(citrus fruit)	(yogurt)	0.082766	0.261671	1.875752

TOP 10 RULES – SORT BY CONFIDENCE

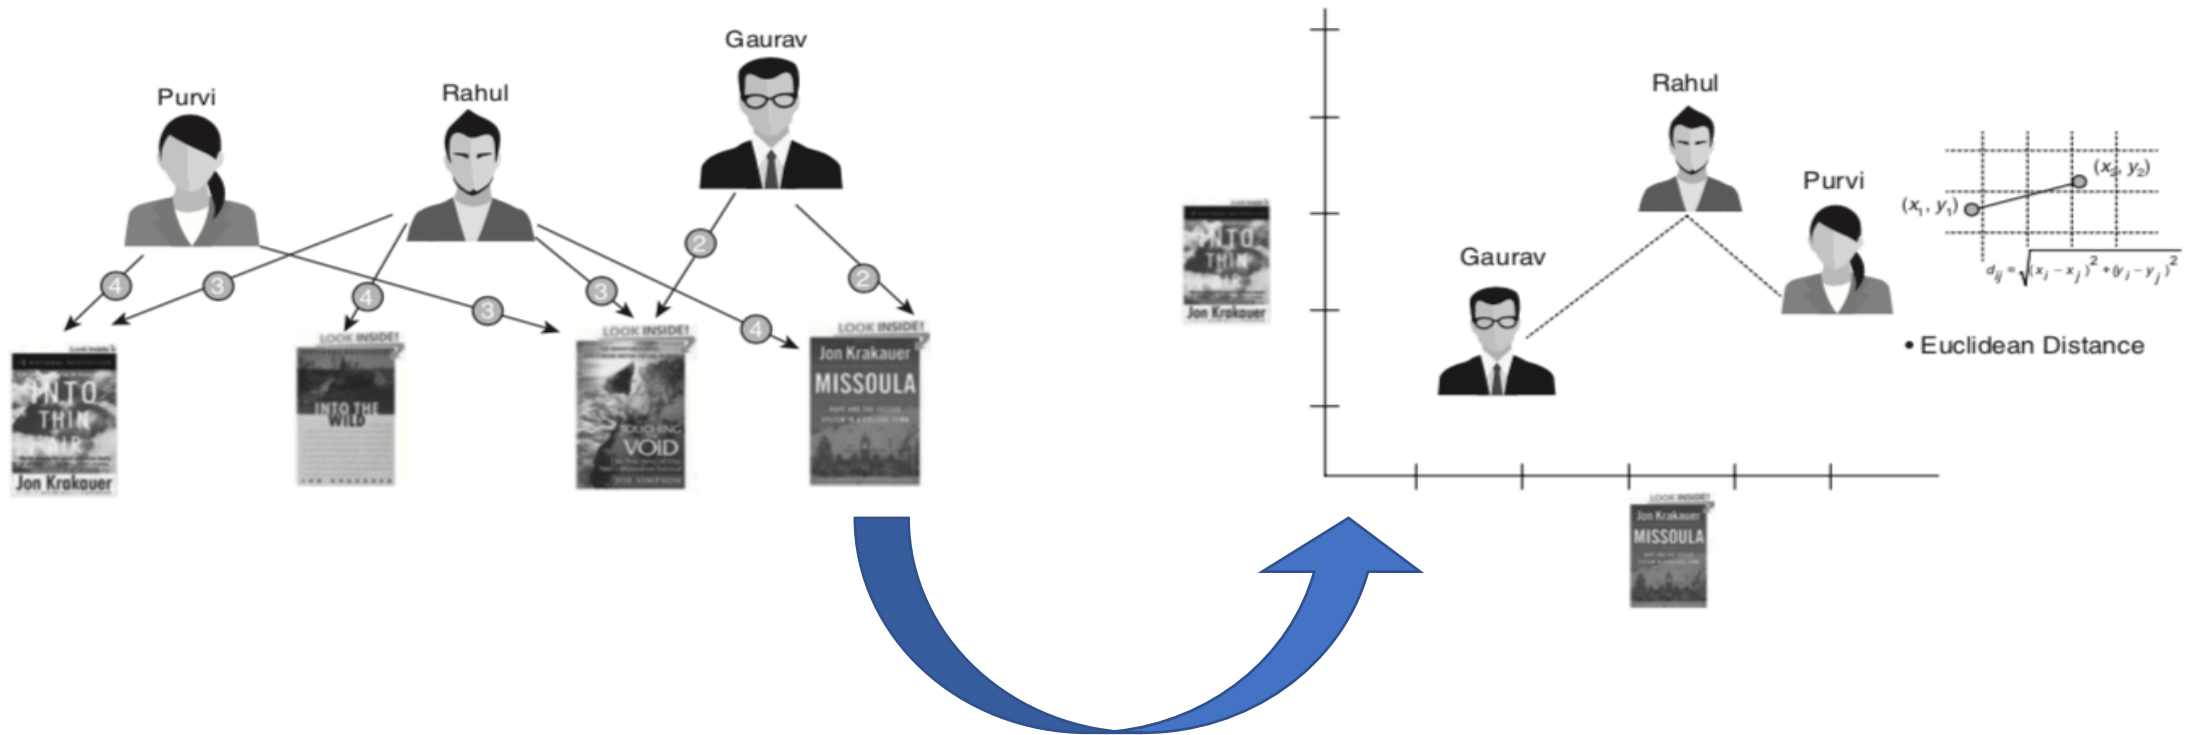
	Antecedants	Consequents	Support	Confidence	Lift
42	(yogurt, other vegetables)	(whole milk)	0.043416	0.512881	2.007235
48	(butter)	(whole milk)	0.055414	0.497248	1.946053
120	(curd)	(whole milk)	0.053279	0.490458	1.919481
80	(other vegetables, root vegetables)	(whole milk)	0.047382	0.489270	1.914833
78	(root vegetables, whole milk)	(other vegetables)	0.048907	0.474012	2.449770
27	(domestic eggs)	(whole milk)	0.063447	0.472756	1.850203
0	(whipped/sour cream)	(whole milk)	0.071683	0.449645	1.759754
89	(root vegetables)	(whole milk)	0.108998	0.448694	1.756031
92	(root vegetables)	(other vegetables)	0.108998	0.434701	2.246605
24	(frozen vegetables)	(whole milk)	0.048094	0.424947	1.663094

PROS AND CONS OF ASSOCIATION RULE MINING

- PROS:
 - Transactions data, which is used for generating rules, is always available and mostly clean
 - The rules generated are simple and can be interpreted
- CONS:
 - Association rules do not take the preference or ratings given by customers into account

COLLABORATIVE FILTERING

How to Find Similarity Between Users?



USER-BASED SIMILARITY

Dataset: <https://grouplens.org/datasets/movielens/>

Features in Dataset:

- userId
- movieId
- Rating
- timestamp

	userId	movieId	rating	timestamp
0	1	31	2.5	1260759144
1	1	1029	3.0	1260759179
2	1	1061	3.0	1260759182
3	1	1129	2.0	1260759185
4	1	1172	4.0	1260759205

USER-BASED SIMILARITY (Contd.)

```
user_movies_df = rating_df.pivot( index='userId', columns='movieId',  
values = "rating" ).reset_index(drop=True)
```

```
user_movies_df.index=rating_df.userId.unique()
```

movieId	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
1	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	4.0	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	4.0	NaN	NaN	NaN	NaN	NaN
5	NaN	NaN	4.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

```
user_movies_df.fillna( 0, inplace = True )
```


COSINE SIMILARITY BETWEEN USERS

```
from sklearn.metrics import pairwise_distances
from scipy.spatial.distance import cosine, correlation

user_sim = 1 - pairwise_distances( user_movies_df.values, metric="cosine" )
user_sim_df = pd.DataFrame( user_sim )

user_sim_df.index = rating_df.userId.unique() user_sim_df.columns =
rating_df.userId.unique()
```

	1	2	3	4	5
1	1.000000	0.000000	0.000000	0.074482	0.016818
2	0.000000	1.000000	0.124295	0.118821	0.103646
3	0.000000	0.124295	1.000000	0.081640	0.151531
4	0.074482	0.118821	0.081640	1.000000	0.130649
5	0.016818	0.103646	0.151531	0.130649	1.000000

	1	2	3	4	5
1	0.000000	0.000000	0.000000	0.074482	0.016818
2	0.000000	0.000000	0.124295	0.118821	0.103646
3	0.000000	0.124295	0.000000	0.081640	0.151531
4	0.074482	0.118821	0.081640	0.000000	0.130649
5	0.016818	0.103646	0.151531	0.130649	0.000000

FILTERING SIMILAR USERS

```
user_sim_df.idxmax(axis=1)[0:5]
```

```
1 325  
2 338  
3 379  
4 518  
5 313
```

```
user_sim_df.iloc[1:2, 330:340]
```

	331	332	333	334	335	336	337	338	339	340
2	0.030344	0.002368	0.052731	0.047094	0.0	0.053044	0.05287	0.581528	0.093863	0.081814

LOADING THE MOVIES DATASET

```
movies_df = pd.read_csv( "ml-latest-small/movies.csv" )
```

	movieid	title	genres
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
1	2	Jumanji (1995)	Adventure Children Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama Romance
4	5	Father of the Bride Part II (1995)	Comedy

COMMON MOVIES OF SIMILAR USERS

```
def get_user_similar_movies( user1, user2 ):
# Inner join between movies watched between two users will give
# the common movies watched.
common_movies = rating_df[rating_df.userId == user1].merge(
rating_df[rating_df.userId == user2], on = "movieId", how = "inner" )
# join the above result set with movies details
return common_movies.merge( movies_df, on = 'movieId' )
```

	userId_x	movieId	rating_x	userId_y	rating_y	title
0	2	17	5.0	338	4.0	Sense and Sensibility (1995)
2	2	47	4.0	338	4.0	Seven (a.k.a. Se7en) (1995)
5	2	150	5.0	338	4.0	Apollo 13 (1995)
28	2	508	4.0	338	4.0	Philadelphia (1993)
29	2	509	4.0	338	4.0	Piano, The (1993)
31	2	527	4.0	338	5.0	Schindler's List (1993)
34	2	589	5.0	338	5.0	Terminator 2: Judgment Day (1991)

USERS WITH DISSIMILAR BEHAVIOR

```
common_movies = get_user_similar_movies( 2, 332 )  
common_movies
```

	userId_x	movieId	rating_x	userId_y	rating_y	title
0	2	552	3.0	332	0.5	Three Musketeers, The (1993)

Challenges with User-Based Similarity

- *cold start* problem in recommender systems

ITEM BASED SIMILARITY

Calculating Cosine Similarity between Movies

```
rating_mat = rating_df.pivot(index='movieId', columns='userId', values =  
"rating").reset_index(drop = True)  
# Fill all NaNs with 0  
rating_mat.fillna(0, inplace = True)  
# Find the correlation between movies  
movie_sim = 1 - pairwise_distances(rating_mat.values,  
metric="correlation") # Fill the diagonal with 0, as it represents the auto-  
correlation # of movies  
movie_sim_df = pd.DataFrame( movie_sim )
```

	0	1	2	3	4
0	1.000000	0.223742	0.183266	0.071055	0.105076
1	0.223742	1.000000	0.123790	0.125014	0.193144
2	0.183266	0.123790	1.000000	0.147771	0.317911
3	0.071055	0.125014	0.147771	1.000000	0.150562
4	0.105076	0.193144	0.317911	0.150562	1.000000

MOST SIMILAR MOVIES

```
def get_similar_movies( movieid, topN = 5 ):
# Get the index of the movie record in movies_df
movieidx = movies_df[movies_df.movieId == movieid].index[0]
movies_df['similarity'] = movie_sim_df.iloc[movieidx]
top_n = movies_df.sort_values( ["similarity"], ascending = False )[0:topN]
return top_n
```

	movieid	title	similarity
695	858	Godfather, The (1972)	1.0

	movieid	title	similarity
695	858	Godfather, The (1972)	1.000000
977	1221	Godfather: Part II, The (1974)	0.709246
969	1213	Goodfellas (1990)	0.509372
951	1193	One Flew Over the Cuckoo's Nest (1975)	0.430101
1744	2194	Untouchables, The (1987)	0.418966

USING SURPRISE LIBRARY

```
from surprise import Dataset, Reader, KNNBasic, evaluate, accuracy

reader = Reader(rating_scale=(1, 5)) data =
Dataset.load_from_df(rating_df[['userId', 'movieId', 'rating']],
reader=reader)
```


USER-BASED SIMILARITY ALGORITHM

```
## Set coefficient and similarity parameters for building model item_based_cosine_sim = {'name': 'pearson',  
'user_based': True}
```

```
knn = KNNBasic(k= 20, min_k = 5, sim_options = item_based_cosine_sim)
```

```
from surprise.model_selection import cross_validate  
cv_results = cross_validate(knn, data, measures=['RMSE'], cv=5, verbose=False)
```

```
np.mean(cv_results.get('test_rmse'))
```

0.9909387452695102

FINDING THE BEST MODEL

```
from surprise.model_selection.search import GridSearchCV
```

```
param_grid = {'k': [10, 20], 'sim_options': {'name': ['cosine', 'pearson'],  
        'User_based': [True, False]}
```

```
grid_cv = GridSearchCV(KNNBasic, param_grid, measures=['rmse'], cv=5,  
        refit=True)
```

```
grid_cv.fit(data)
```

```
print(grid_cv.best_params['rmse'])
```

0.9963783863084851

	param_k	param_sim_options	mean_test_rmse	rank_test_rmse
0	10	{'name': 'cosine', 'user_based': True}	1.009724	4
1	20	{'name': 'cosine', 'user_based': True}	0.996378	1
2	10	{'name': 'cosine', 'user_based': False}	1.048802	8
3	20	{'name': 'cosine', 'user_based': False}	1.015225	6
4	10	{'name': 'pearson', 'user_based': True}	1.012283	5
5	20	{'name': 'pearson', 'user_based': True}	1.000766	2
6	10	{'name': 'pearson', 'user_based': False}	1.030900	7
7	20	{'name': 'pearson', 'user_based': False}	1.004205	3

MATRIX FACTORIZATION

Users–Movies Rating Matrix

		Movies				
		M1	M2	M3	M4	M5
Users	U1	3	4	2	5	1
	U2	2	4	1	2	4
	U3	3	3	5	2	2

Users–Factors Matrix

		Factors		
		F1	F2	F3
Users	U1	0.73	3.22	0
	U2	0	1.57	2.53
	U3	1.62	0	1.44

Factors–Movies Matrix

		Movies				
		M1	M2	M3	M4	M5
Factors	F1	1.47	1	2.73	1.73	0
	F2	0.6	1.01	0	1.27	0.31
	F3	0.42	0.95	0.39	0	1.39

MATRIX FACTORIZATION (Contd.)

```
from surprise import SVD
```

```
# Use 10 factors for building the model
```

```
svd = SVD( n_factors = 5 )
```

```
cv_results = cross_validate(svd, data, measures=['RMSE'], cv=5,  
verbose=True)
```

```
Evaluating RMSE of algorithm SVD on 5 split(s).
```

	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Mean	Std
RMSE (testset)	0.8904	0.8917	0.8983	0.8799	0.8926	0.8906	0.0060
Fit time	1.91	1.99	1.96	1.87	1.84	1.91	0.06
Test time	0.21	0.18	0.18	0.17	0.18	0.18	0.01

Thank You!